

For this part of assignment, let us use the module 3 movie ratings data and attempt to predict missing ratings using non-negative matrix factorization.

Section 1

```
In [16]: import pandas as pd
pd.set_option('display.max_columns', None)
```

```
# import data and display columns
data_users = pd.read_csv('users.csv')
data_movies = pd.read_csv('movies.csv')
data_train = pd.read_csv('train.csv')
data_test = pd.read_csv('test.csv')
print(data_users.head())
print(data_movies.head())
print(data_train.head())
print(data_test.head())
```

	uID	gender	age	occupation	zip
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

	mID		title	year	Doc	Com	Hor	Adv	Wes	Dra	Ani	\
0	1		Toy Story	1995	0	1	0	0	0	0	1	
1	2		Jumanji	1995	0	0	0	1	0	0	0	
2	3		Grumpier Old Men	1995	0	1	0	0	0	0	0	
3	4		Waiting to Exhale	1995	0	1	0	0	0	1	0	
4	5		Father of the Bride Part II	1995	0	1	0	0	0	0	0	

	War	Chi	Cri	Thr	Sci	Mys	Rom	Fil	Fan	Act	Mus
0	0	1	0	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0	1	0	0
2	0	0	0	0	0	0	1	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	0

	uID	mID	rating
0	744	1210	5
1	3040	1584	4
2	1451	1293	5
3	5455	3176	2
4	2507	3074	5

	uID	mID	rating
0	2233	440	4
1	4274	587	5
2	2498	454	3
3	2868	2336	5
4	1636	2686	5

```
In [17]: import numpy as np
from sklearn.decomposition import NMF
from sklearn.metrics import mean_squared_error

# create user-movie matrix:
# rows: users
# cols: movies
# vals: ratings
mat_user_movie = data_train.pivot(index='uID', columns='mID', values='rating').fillna(0)

# use NMF
mod_nmf = NMF(n_components=15, max_iter=1500, random_state=42)
feat_user = mod_nmf.fit_transform(mat_user_movie)
feat_movie = mod_nmf.components_

# create predicted ratings matrix
mat_predict_ratings = np.dot(feat_user, feat_movie)

# get predicted ratings for test data
# check if movie ID exists in training data and default to 0 if it doesn't
test_uIDs = data_test['uID']
test_mIDs = data_test['mID']
train_uIDs = mat_user_movie.index
train_mIDs = mat_user_movie.columns
test_ratings = data_test['rating']
test_predicts = []
for uID, mID in zip(test_uIDs, test_mIDs):
    if (uID in train_uIDs) and (mID in train_mIDs):
        uID = train_uIDs.get_loc(uID)
        mID = train_mIDs.get_loc(mID)
        test_predicts.append(mat_predict_ratings[uID, mID])
```

```
        else:
            test_predicts.append(0)

# get RMSE
print(f'RMSE: {np.sqrt(mean_squared_error(test_ratings, test_predicts))}')
```

RMSE: 2.8724517237578873

Section 2

So, it seems we ended up with a rather high RMSE of ~2.87. This is worse than what the simpler baseline and similarity-based methods yielded and shows that NMF might not be a good choice for this kind of problem. In fact, there are several reasons why this might be the case.

To start with, NMF requires a rather dense matrix to be accurate. After all, it works by learning latent factors. Our matrices though, were quite sparse. This also reflects the real world where people do not usually go through a long list of movies, leaving ratings or feedback for each. Furthermore, NMF will always struggle with users/movies that have few ratings, which is very common with new users and movies. NMF is also highly susceptible to overfitting—it might try to learn too many latent factors from the training data and fail to apply that knowledge to the test data.

Some ways to potentially improve an NMF-based predictor is to leverage regularization. This might help with overfitting and allow for better generalization when working with test data. Furthermore, we could incorporate some similarity-based methods into our NMF model, creating a hybrid system. Such a system would likely use NMF to extract the latent factors but then rely on a collaborative approach for final predictions. Another way of helping NMF perform better is to give it more data in addition to ratings, such as whether a user finished the movie or how many times they watched it.

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js